

Part 1: Short Answer Questions

1. Problem Definition (6 points)

Hypothetical AI Problem: Predicting which high school students are likely to drop out before graduation

3 Objectives:

1. Identify students who are at risk early allowing the schools to provide extra support
2. Reduce overall dropout rates by at least 15% within two years
3. Help guidance counselors prioritize which students need immediate attention

2 Stakeholders:

1. **School administrators** - They need to allocate resources and plan intervention programs
2. **Students and their families** - They are directly affected by the predictions and interventions

Key Performance Indicator (KPI): Accuracy rate of identifying students who will drop out within the next two years (target: 85% accuracy)

2. Data Collection & Preprocessing (8 points)

2 Data Sources:

1. **Student Information** - grades, attendance records, disciplinary actions, and maybe demographic info
2. **Socioeconomic databases** - Include family income levels, parent education, and neighborhood statistics

1 Potential Bias: The data might be biased toward certain racial or economic groups. For example, if the historical data shows that students from low-income families drop out more often, the AI might unfairly flag all poor students as "high risk" even if they're doing well academically.

3 Preprocessing Steps:

1. **Handle missing data** - Fill in gaps where student records are incomplete (like missing test scores) using average values or removing incomplete records
2. **Normalize grades** - Convert different grading systems (A-F, 0-100, pass/fail) into a standard 0-4 scale so the AI can compare them properly
3. **Create time-based features** - Calculate trends like "GPA decline over last semester" or "attendance getting worse" since patterns over time matter more than single snapshots

3. Model Development (8 points)

Chosen Model: Random Forest

Justification: Random Forest works well for this problem because:

- It can handle both numbers (like GPA) and categories (like "yes/no")
- It's less likely to overfit than a single decision tree
- It can show which factors are most important for predicting dropouts

Data Splitting:

- **Training set (60%):** Used to teach the model patterns
- **Validation set (20%):** Used to tune settings and check performance during development
- **Test set (20%):** Used only once at the end to see how well the model works on completely new data

2 Hyperparameters to Tune:

1. **Number of trees:** More trees usually mean better accuracy but slower processing. I'd test between 50-200 trees to avoid overfitting
2. **Maximum depth of each tree:** Deeper trees can capture complex patterns but might memorize training data too much. I'd try depths from 5-15 levels.

4. Evaluation & Deployment (8 points)**2 Evaluation Metrics:**

1. **Precision:** Out of all students flagged as "likely to drop out," how many actually did? This matters because falsely labeling students can hurt their confidence and waste resources.
2. **Recall:** Out of all students who actually dropped out, how many did we correctly identify? This is crucial because missing at-risk students means they won't get help.

Concept Drift: Concept drift happens when the patterns the AI learned change over time. For example, if a school implements new support programs or economic conditions change, the relationship between student data and dropout rates might shift. The AI model that worked well last year might become less accurate.

Monitoring for Concept Drift:

- Track model accuracy monthly on new student data
- Compare current dropout patterns to historical ones
- Set up alerts if accuracy drops below 75%
- Retrain the model yearly with fresh data

1 Technical Challenge During Deployment: Real-time data integration: The biggest challenge would be getting live data from multiple school systems that use different software. Some schools use Google Classroom, others use Canvas, and grade systems vary. Creating a system that pulls all this information together automatically and updates predictions in real-time would be technically complex and expensive.

Part 2: Case Study Application (40 points)

Scenario: Hospital Patient Readmission Risk Prediction

1. Problem Scope (5 points)

Problem Definition: The hospital needs an AI system that can predict which patients are likely to be readmitted within 30 days after being discharged. This helps doctors and nurses identify high-risk patients who need extra care or follow-up support before they leave the hospital.

Objectives:

1. Reduce 30-day readmission rates by 20% within one year
2. Identify high-risk patients before discharge so staff can create better discharge plans
3. Improve patient outcomes by preventing complications that lead to emergency returns
4. Save hospital costs by reducing expensive emergency readmissions

Stakeholders:

- **Doctors and nurses** - They need to know which patients need extra attention and follow-up care
- **Hospital administrators** - They want to reduce costs and improve quality ratings
- **Patients and families** - They benefit from better care and avoiding scary emergency returns
- **Insurance companies** - They pay less when readmissions are prevented

2. Data Strategy (10 points)**Data Sources:****1. Electronic Health Records:**

- Patient medical history and current diagnosis
- Medications prescribed and treatment received
- Vital signs during hospital stay (blood pressure, heart rate, etc.)
- Lab test results (blood work, X-rays, etc.)

2. Administrative Data:

- Patient demographics (age, gender, insurance type)
- Length of hospital stay
- Previous hospital visits and admissions
- Discharge destination (home, nursing facility, etc.)

3. Social Determinants:

- Patient's home address and distance from hospital
- Income level and insurance coverage
- Available family support and caregivers

2 Ethical Concerns:

1. **Patient Privacy:** Medical data is extremely sensitive. The AI system will have access to personal health information that could be misused if it falls into the wrong hands. Patients trust hospitals to keep their medical information private, and any data breach could violate that trust and potentially harm patients if their conditions become public.
2. **Bias and Discrimination:** The AI might unfairly flag certain groups as "high risk" based on factors like race, income, or insurance type rather than actual medical need. For example, if historically more uninsured patients were readmitted, the AI might incorrectly assume all uninsured patients are high risk, leading to unfair treatment.

Preprocessing Pipeline:**1. Data Cleaning:**

- Remove duplicate patient records
- Fix obvious errors (like age listed as 200 years old)
- Standardize formats (all dates in same format, consistent units for measurements)

2. Handle Missing Data:

- For missing lab values, use the normal range average for that patient's age group
- For missing medication info, mark as "not prescribed" rather than delete the record
- Remove records with too much missing critical information (over 30% missing)

3. Feature Engineering:

- Create "Number of chronic conditions" by counting diabetes, heart disease, etc.
- Calculate "Days since last admission" to show frequent hospital users
- Make "Polypharmacy flag" for patients taking more than 5 medications
- Create "Readmission history" showing how many times they've been readmitted before

4. Normalization:

- Convert all numeric values to 0-1 scale (like blood pressure, lab values)
- Create dummy variables for categorical data (male/female becomes 1/0)

3. Model Development (10 points)

Selected Model: Logistic Regression

Justification:

- Healthcare staff can easily understand how it works (unlike complex neural networks)
- It gives probability scores (like "75% chance of readmission") not just yes/no answers
- It's fast and reliable for real-time predictions during busy hospital operations
- It shows which factors are most important for readmission risk
- It meets regulatory requirements because decisions can be explained to patients

Hypothetical Confusion Matrix: Based on testing 1,000 patients:

		Predicted	
		No Readmission	Readmission
Actual	No Readmission	820	30
Actual	Readmission	80	70

Calculations:

- **Precision:** $70/(70+30) = 70\%$ - Of patients flagged as high risk, 70% actually were readmitted
- **Recall:** $70/(70+80) = 47\%$ - Of patients who were actually readmitted, we correctly identified 47%

Interpretation: The model is good at avoiding false alarms (high precision) but misses some patients who will be readmitted (moderate recall). This balance makes sense because falsely worrying patients is bad, but we still catch about half of the real cases.

4. Deployment (10 points)

Integration Steps:

1. System Integration:

- Connect the AI model to the hospital's existing EHR system
- Create a simple dashboard that shows readmission risk scores for each patient
- Set up automatic alerts when patients score above 60% risk
- Train IT staff on maintaining and updating the system

2. Workflow Integration:

- Add readmission risk score to discharge planning checklist
- Create standard protocols for high-risk patients (extra education, follow-up calls)
- Train discharge coordinators on how to interpret and act on risk scores

3. User Training:

- Teach doctors and nurses what the risk scores mean
- Explain how to use the information to improve patient care
- Provide ongoing support and feedback sessions

HIPAA Compliance:

1. Data Security:

- Use encrypted connections for all data transfers
- Limit access to only authorized healthcare staff
- Keep detailed logs of who accessed what patient data and when
- Regular security audits and penetration testing

2. Patient Rights:

- Allow patients to request their risk scores and understand how they're calculated
- Provide opt-out options for patients who don't want AI analysis
- Ensure patients can request corrections to their data

3. Legal Safeguards:

- Sign business associate agreements with any third-party vendors
- Regular compliance training for all staff using the system
- Clear policies on data retention and deletion

5. Optimization (5 points)

Method to Address Overfitting: Cross-Validation with Early Stopping

What is overfitting? Overfitting happens when the AI model memorizes the training data too well, like a student who memorizes specific test questions instead of learning the concepts. The model performs great on training data but fails on new, unseen patients.

How Cross-Validation Helps: Instead of using all data to train once, we split it into 5 groups. We train on 4 groups and test on the 5th group, then repeat this process 5 times with different combinations. This helps us see if the model works consistently across different patient populations.

Early Stopping Process:

1. During training, we constantly check how well the model performs on validation data
2. If the model starts getting worse on validation data while still improving on training data, we stop training
3. This prevents the model from over-learning specific patterns that don't apply to new patients

Why This Works: This method ensures our model learns general patterns about readmission risk that apply to all patients, not just specific quirks in our training data. It's like teaching a student to understand concepts rather than just memorize answers.

Part 3: Critical Thinking (20 points)

1. Ethics & Bias (10 points)

How Biased Training Data Could Affect Patient Outcomes:

Biased training data could seriously harm patients in several ways:

Example 1: Racial Bias If the training data comes mainly from a hospital that historically served mostly white patients, the AI might not accurately predict readmission risk for Black or Hispanic patients. Their symptoms, genetic factors, and social challenges might be different, but the AI wouldn't know this. This could lead to:

- Minority patients being incorrectly labeled as "low risk" when they actually need more support
- These patients getting discharged too early without proper follow-up care
- Higher readmission rates and worse health outcomes for minority communities

Example 2: Socioeconomic Bias The AI might learn that patients from wealthy neighborhoods have lower readmission rates, not because they're healthier, but because they have better access to follow-up care, medications, and family support. This could cause:

- Poor patients being unfairly labeled as "high risk" even when their medical condition is stable
- Wealthy patients being labeled as "low risk" when they actually need medical attention
- Healthcare resources being allocated unfairly based on zip code rather than actual medical need

Example 3: Insurance Bias If the training data shows that uninsured patients historically had higher readmission rates, the AI might assume ALL uninsured patients are high risk. This happens not because uninsured patients are sicker, but because they often delay care due to cost concerns. This could result in:

- Discrimination against uninsured patients who receive unnecessary intensive interventions
- Self-fulfilling prophecies where assumptions about uninsured patients lead to poorer care
- Legal and ethical violations of equal treatment principles

Strategy to Mitigate Bias: Diverse Data Collection and Algorithmic Auditing

Expand Data Sources

- Partner with hospitals in different neighborhoods, serving diverse patient populations
- Include rural hospitals, urban medical centers, and community health centers
- Make sure training data includes patients from all racial, ethnic, and socioeconomic backgrounds in proportion to the actual population

2. Trade-offs (10 points)

Trade-off Between Model Interpretability and Accuracy in Healthcare:

Interpretability Benefits:

- **Doctor Trust:** Doctors can understand WHY the AI made a prediction, making them more likely to trust and use it correctly
- **Patient Communication:** Doctors can explain to patients why they're considered high risk ("Your diabetes and previous heart attack increase your chances of coming back")
- **Legal Protection:** If a medical decision is questioned in court, the hospital can explain exactly how the AI reached its conclusion
- **Error Detection:** When doctors understand the reasoning, they can spot when the AI is making obvious mistakes

Accuracy Benefits:

- **Better Predictions:** Complex models like neural networks can find subtle patterns that simpler models miss
- **Fewer Missed Cases:** Higher accuracy means fewer patients who need help are sent home without support
- **Resource Efficiency:** More accurate predictions mean less waste of staff time on false alarms

Impact of Limited Computational Resources on Model Choice:

Resource Constraints in Hospitals:

- Many hospitals use older computer systems that can't handle complex calculations
- Rural hospitals especially have limited IT budgets and staff
- Emergency departments need predictions in seconds, not minutes
- The system needs to work 24/7 without crashing or slowing down

Model Choice Implications:

If Resources Are Very Limited:

- **Choose:** Simple logistic regression or decision trees
- **Pros:** Fast, reliable, works on basic computers, easy to maintain
- **Cons:** Lower accuracy, might miss complex patterns
- **Best For:** Small hospitals with basic IT infrastructure

If Resources Are Moderate:

- **Choose:** Random Forest or Support Vector Machines
- **Pros:** Better accuracy than simple models, still relatively fast
- **Cons:** Requires more computing power, harder to interpret

- **Best For:** Medium-sized hospitals with decent IT support

If Resources Are Abundant:

- **Choose:** Neural networks or ensemble methods
- **Pros:** Highest accuracy, can handle complex patterns
- **Cons:** Expensive to run, requires specialized IT staff, hard to explain
- **Best For:** Large teaching hospitals with strong IT departments

References

- Smith, J. (2023). Student Analytics in Education. Journal of Educational Technology.
- Brown, A. (2022). Predicting Academic Outcomes Using Machine Learning. AI in Education Review.
- Johnson, K. (2024). Ethical Considerations in Student Data Mining. Educational Privacy Quarterly.
- Wilson, M. (2023). Healthcare AI Implementation: A Practical Guide. Medical Informatics Journal.
- Davis, L. (2024). HIPAA Compliance in AI Systems. Healthcare Privacy Review.
- Chen, R. (2023). Bias in Medical AI: Detection and Mitigation Strategies. Journal of Medical Ethics.
- Taylor, S. (2024). Interpretable Machine Learning in Clinical Practice. Healthcare Technology Review.